


Figure F1. Institutional preference depends on type and ex ante exposure

Outside-option slice $\phi_0 = 0.250 \times \phi_1$; $m = 4$ weak states; $\beta = 0.90$



Colored polygons report the payoff comparison between unanimity and majority. The first two panels condition on H's type. The ex ante panel applies the prior weights $(1-\nu, \nu)$ to every type-payoff vector passed through the frozen N6 interface before taking unanimity minus majority; it adds no equilibrium selection or public benchmark. The neutral hatched region has no comparison because unanimity has no pure-vote PBE. The left-edge segments report the distinct complete-information endpoint at $\nu = 0$. Boundaries are $\nu^* = (1-\kappa)\phi_1/(1-\kappa\phi_1)$, $\nu_{SP} = \beta(1-\kappa)\phi_1/[1-\beta\kappa\phi_1-\beta(q-1)/m]$, $\nu_{SE} = \beta(1/m-\kappa\phi_1)/[\beta(1/m-\kappa\phi_1)+1-\beta q/m]$, and $\nu_{XA} = (\beta-\kappa)/(1-\kappa)$ on the majority-exclusion branch. The vertical ν_{XA} line is specific to the displayed slice $\phi_0 = \kappa\phi_1$, where ϕ_1 cancels; if ϕ_0 were held fixed, the ex ante boundary would vary with ϕ_1 . Because $b_\theta = 0$ by construction, the strong type has no direct membership benefit and its only gain from inclusion is the concession; its preference for exclusion reflects the baseline's isolation of informational power, not a general substantive prediction. The dashed horizontal line is the substitute cost $\phi_1 = 1/m$. Note: Model-generated regions from the frozen private-information payoff vectors; all boundaries closed-form.