


Figure F1. Institutional preference depends on type and ex ante exposure

Outside-option slice $\phi_0 = 0.500 \times \phi_1$; $m = 4$ weak states; $\beta = 0.90$



Analytical boundary -- ν_{SE} --- ν_{SP} ν_{XA} — ν_{star}

H's payoff comparison (private information) H indifferent H prefers majority H prefers unanimity No comparison: no pure-vote PBE (unanimity)

Colored polygons report the payoff comparison between unanimity and majority. The first two panels condition on H's type. The ex ante panel applies the prior weights $(1 - \nu, \nu)$ to every type-payoff vector passed through the frozen N6 interface before taking unanimity minus majority; it adds no equilibrium selection or public benchmark. The neutral hatched region has no comparison because unanimity has no pure-vote PBE. The left-edge segments report the distinct complete-information endpoint at $\nu = 0$. Boundaries are $\nu_{\text{star}} = (1 - \kappa)\phi_1 / (1 - \kappa\phi_1)$, $\nu_{\text{SP}} = \beta(1 - \kappa)\phi_1 / [1 - \beta\kappa\phi_1 - \beta(q - 1)/m]$, $\nu_{\text{SE}} = \beta(1/m - \kappa\phi_1) / [\beta(1/m - \kappa\phi_1) + 1 - \beta q/m]$, and $\nu_{\text{XA}} = (\beta - \kappa) / (1 - \kappa)$ on the majority-exclusion branch. The dashed horizontal line is the hegemony threshold $m \times \phi_1 = 1$. Note: Model-generated regions from the frozen private-information payoff vectors; all boundaries closed-form.