


Figure F1. Institutional preference depends on type and ex ante exposure

Outside-option slice $o_0 = 0.750 \times o_1$; $m = 4$ weak states; $\beta = 0.90$



Analytical boundary -- nu_SE --- nu_SP nu_XA — nu_star

H's payoff comparison (private information) H indifferent H prefers majority H prefers unanimity No comparison: no pure-vote PBE (unanimity)

Colored polygons report the payoff comparison between unanimity and majority. The first two panels condition on H's type. The ex ante panel applies the prior weights $(1-\nu, \nu)$ to every type-payoff vector passed through the frozen N6 interface before taking unanimity minus majority; it adds no equilibrium selection or public benchmark. The neutral hatched region has no comparison because unanimity has no pure-vote PBE. The left-edge segments report the distinct complete-information endpoint at $\nu = 0$. Boundaries are $\nu^* = (1-\kappa)o_1/(1-\kappa o_1)$, $\nu_{SP} = \beta(1-\kappa)o_1/[1-\beta\kappa o_1-\beta(q-1)/m]$, $\nu_{SE} = \beta(1/m-\kappa o_1)/[\beta(1/m-\kappa o_1)+1-\beta q/m]$, and $\nu_{XA} = (\beta-\kappa)/(1-\kappa)$ on the majority-exclusion branch. The dashed horizontal line is the substitute cost $o_1 = 1/m$. Note: Model-generated regions from the frozen private-information payoff vectors; all boundaries closed-form.