

Informational Power Through Pivotality: How Consensus Can Benefit a Hegemon

Manoel Galdino (University of São Paulo; mgaldino@usp.br)

2026-08-22

Abstract

Proposal rights are a well-documented source of political power. This paper asks whether a hegemon that holds none retains another: private information, converted into concessions by the voting rule. In a two-round bargaining game, only weak states propose, and a hegemon privately knows whether its disagreement payoff is low or high; majority and unanimity differ solely in the number of votes required. Under majority, a proposer can replace the hegemon with uninformed weak-state votes, and this substitute caps the value of its private information. Under unanimity, the hegemon's vote is an essential input and proposals must respond to a threshold weak states cannot observe. To separate this informational component from the value of a necessary vote, we benchmark each private game against its complete-information counterpart and compare the resulting type-specific informational rents across rules. The contrast is conditional. When the high type is sufficiently likely, it is positive for the low type whenever majority screens or excludes the hegemon with or without information, zero when majority already pays the pooling price, and of either sign when private majority excludes a low type that public majority would include; the high type gains only when private majority screens where public majority would exclude it. For intermediate beliefs, unanimity has no perfect Bayesian equilibrium in pure ballot strategies, so the comparison is undefined there. Private information can thus substitute for agenda control as a source of hegemonic power, but only where the voting rule removes substitutes for the informed vote.

1 Introduction

At the height of American power, immediately after the end of the Cold War, the United States led the creation of the World Trade Organization. Under its rules, states hold equal voting power, the United States took no agenda power, and agreements require consensus.¹

The WTO is, of course, no exception in how international organizations set their rules (Koremenos, Lipson, and Snidal 2001). What is puzzling is the distance between formal equality among states and the asymmetry of power at the height of American primacy in a unipolar order, against what any realist account would lead one to expect (see, for example, Krasner 1983). The distance is widest on the greatest source of power of all, agenda power (Kalandrakis 2006).

The explanation offered in the literature is that, despite formal equality, some form of informal inequality must exist. Steinberg (2002) documents, for the WTO, agenda power confined to a handful of countries through devices such as the green room, alongside other mechanisms through which the United States could steer the organization in a direction of its own choosing.

The United States thus holds neither exclusive proposal rights nor extra votes, and it knows more than its partners about what it can obtain if no agreement is reached. Where does its influence come from?

The standard answer is proposal rights. Kalandrakis (2006) shows that the right to decide what is put to a vote translates into political power, and legislative-bargaining models since Baron and Ferejohn (1989) have traced how recognition rules and agenda control distribute surplus (Winter 1996; Eraslan and Evdokimov 2019). This paper removes that source deliberately. The hegemon never proposes, and its vote counts as one among many. The question then becomes whether any other source of power remains. The answer we develop is that private information can be one, but only under voting rules that make the informed vote impossible to replace.

The institutional version of the question is when unanimity advantages a hegemon over majority rule, given equal votes and no hegemonic agenda control, and how much of that advantage arises from private information rather than from the veto that unanimity confers on everyone. The second half matters because a necessary vote is valuable even when its holder's strength is common knowl-

¹For the model, *consensus* and *unanimity* denote the same voting rule: every state must approve. The distinction between the terms does not affect the analysis.

edge. The informational puzzle begins where the complete-information benchmark ends: what does a hegemon gain because the other states cannot price its vote?

We study a two-round bargaining game with one hegemon and at least three weak states. Only weak states propose. The hegemon privately knows whether its terminal disagreement payoff is low or high. Proposals and simultaneous ballots are public, the institutional surplus is fixed, and a failed first round leads to a discounted terminal round. Majority and unanimity differ only in the number of votes required.

The mechanism is coalition substitutability. Under majority, a weak-state proposer can assemble a winning coalition from uninformed weak-state votes and leave the hegemon out. That outside coalition caps the price of the hegemon's vote, and with it the value of anything the hegemon privately knows. Under unanimity the substitute disappears. Every agreement requires the hegemon's approval, so the proposal must respond to a threshold that weak states cannot observe. When the high type is sufficiently likely, the proposer pays the high threshold to both types, and a low-type hegemon collects the concession designed for a stronger one. The rent arises from pooling: weak states price the hegemon as strong because a failed proposal has no fallback coalition.

Comparing unanimity with majority does not by itself identify this rent. A hegemon may receive more under unanimity simply because its vote has become necessary, an advantage that exists under complete information. We therefore solve each voting rule twice, with the hegemon's type public and private, and define the informational rent of a rule as the difference between its private and public payoffs for each type. The difference of these rents across rules measures what unanimity adds to the return on private information once the value of a necessary vote has been removed.

That measure is conditional in ways a single-number comparison would hide. Majority has pure-strategy equilibria featuring exclusion, type-contingent screening, or pooling, depending on outside options and beliefs. Unanimity has an immediate low-type agreement when the low type is certain, an immediate pooling agreement when the high type is sufficiently likely, and no perfect Bayesian equilibrium in pure ballot strategies between these cases. Above the pooling cutoff, the low type's informational advantage under unanimity is positive when majority screens or when majority excludes the hegemon with or without information, zero when majority already pools, and of either sign when private majority excludes a low type that public majority would include. The high type

gains only when private majority screens where public majority would exclude it, and loses when private majority excludes a high type that public majority would include. Where private unanimity has no pure-strategy equilibrium the contrast is empty; we report that emptiness as a boundary of the equilibrium domain and do not fill it with a counterfactual payoff.

Two recent papers also make disagreement values private and compare majority with unanimity, through distinct mechanisms. Piazzolo and Vanberg (2025) study dynamic signaling by privately informed responders: under unanimity, rejection is a valuable signal of toughness, whereas majority disciplines such signaling through the threat of exclusion, so votes are more expensive under unanimity. Glynnia, Thum, and Xefteris (2026) shut down the dynamic channel and show that uncertainty about several responders changes the proposer’s choice between insuring against costly types and gambling on cheap ones. In both papers, private information is dispersed among symmetric responders and the object of analysis is the price of votes or the probability of agreement. Miller, Montero, and Vanberg (2018) supply the complete-information benchmark with heterogeneous disagreement values, in which an expensive player can be included under unanimity and excluded under majority.

We ask a distributional question about one asymmetric actor, and three features of the model separate it from this work. Private information is concentrated in a hegemon that never proposes, so any rent it earns cannot be attributed to agenda power. The voting rule determines whether other votes substitute for the informed one, and substitutability determines whether private information can be converted into concessions at all. And the complete-information benchmark for each rule separates the informational component of unanimity’s advantage from its veto component. The threat of exclusion in Piazzolo and Vanberg (2025) and the coalition substitute here can be read as the same institutional fact from two sides: they trace its consequences for responders’ incentives to signal, we trace its consequences for the rents of the actor whose vote the rule makes replaceable or essential. Rational-design and informal-power accounts of international institutions (Koremenos, Lipson, and Snidal 2001; Steinberg 2002; Stone 2011) document influence that formal rules do not record; the model specifies one channel for that influence and the voting rules under which it operates.

Section 3 defines the game and solution concept. Section 4 solves the complete-information benchmarks, then the private games, and finally the informational rents and their difference across voting

rules. Section 5 interprets the essential-input mechanism and its limits. The appendix gives the proofs, endpoint discipline, exact correspondence tables, and a worked numerical illustration.

2 A working numerical illustration

Consider one hegemon and four weak states. The unit pie is bargained over with discount factor 0.9. The two disagreement payoffs are 0.10 and 0.35. In the terminal unanimity round, a proposer chooses between an offer accepted only by the low type and an offer accepted by both types. The cutoff probability of the high type is

$$\nu^* = \frac{0.35 - 0.10}{1 - 0.10} = 0.2778.$$

At a belief of 0.20, the proposer offers 0.10 and accepts the risk of failure when the type is high. At a belief of 0.35, it offers 0.35 and both types accept. This is a working numerical illustration, not an empirical calibration. The full two-round game adds the option to buy weak-state votes and the continuation created by a failed first ballot.

3 The model

3.1 Players, information, and proposals

There is one hegemon, H , and $m \geq 3$ weak states, $W = \{1, \dots, m\}$. Thus $N = m + 1$. Nature selects $\theta \in \{0, 1\}$, observed only by H , with $\Pr(\theta = 1) = \nu \in [0, 1]$. The terminal disagreement payoffs satisfy

$$0 < o_0 < o_1 < 1.$$

The game has two rounds. In each round, one weak state is recognized uniformly to propose; H is never a proposer. The two recognition draws are independent and made with replacement. Every weak state remains eligible in Round 2, including the Round-1 proposer. A proposal by weak state

i is

$$s = (y, (x_j)_{j \in W \setminus \{i\}}, r_i), \quad 0 \leq y \leq \bar{y}, \quad x_j \geq 0, \quad r_i \geq 0, \quad y + \sum_{j \neq i} x_j + r_i \leq 1,$$

where $o_1 \leq \bar{y} \leq 1$. Payments cannot be negative and side payments outside the proposed package are unavailable. Here y is the institutional concession assigned to H , x_j is the share assigned to responding weak state j , and r_i is the proposer's residual. The fixed unit pie is exhausted on the equilibrium path.

The payoff from agreement for H is $y + b_\theta$. We normalize $b_\theta = 0$, so an included hegemon receives the concession y . [AUTHOR: P2] This normalization isolates informational concessions and pivotality from any intrinsic benefit that an agreement might itself provide.

3.2 Ballots, timing, and payoffs

The proposer is counted as voting yes. All other states vote simultaneously; the complete vote vector and outcome become public only after all votes have been cast. Majority requires

$$q = \left\lfloor \frac{N}{2} \right\rfloor + 1$$

yes votes, while unanimity requires every state's yes vote. The ballot protocol is otherwise identical across rules.

If a proposal passes, every stated weak-state allocation is implemented. If H votes yes, it receives y . If H votes no but a majority passes the proposal without it, its general payoff is $y + o_\theta$. This formula is part of the complete game because the proposed institutional term y is implemented even without H 's approval. In every equilibrium exclusion derived below, however, the proposer chooses $y = 0$; the hegemon therefore receives exactly o_θ , never a positive concession on top of it.

If nothing passes in Round 1, no payoff is realized then and the game proceeds to Round 2. Round-2 payoffs are multiplied by $\beta \in (0, 1)$ when evaluated in Round-1 units. If nothing passes in Round 2, each weak state receives zero and H receives o_θ in Round-2 units. A no vote is only a ballot action; it does not remove a player or trigger disagreement by itself. [AUTHOR: P1] The delayed

terminal disagreement payoff represents the cost of prolonging international negotiations.

Table 1: Transitions and payoffs in the essential-input game

Event	Weak-state payoffs	H votes yes	H votes no
Proposal passes	Proposed x_j and r_i	y	$y + o_\theta$
Round-1 proposal fails	No current payoff; proceed to Round 2	$\beta C_H(h^Y)$	$\beta C_H(h^N)$
Round-2 proposal fails	Zero	o_θ	o_θ

Here h^Y and h^N denote the distinct public histories in which H voted yes or no, and $C_H(h^a)$ denotes H 's Round-2 continuation value after history h^a . These continuation values need not be equal because the public vote can change beliefs before Round 2.




Figure 1: Sequence of proposals and ballots. Within each ballot, responding states vote simultaneously. A failed first-round proposal continues to the terminal round; a no vote alone does not end the game.

3.3 Solution concept

We use perfect Bayesian equilibrium with pure ballot strategies and three declared disciplines. First, an action by a player who does not know θ does not change beliefs about θ . After an action by H , beliefs follow Bayes' rule whenever possible and otherwise satisfy structural consistency. Second, a weak responder evaluates yes and no as if its own vote were pivotal. Third, exact expected indifference is resolved in favor of yes, denoted T^Y . Among proposal choices that give the proposer the same expected payoff, the selected proposal minimizes H 's expected payoff.

Endpoint beliefs preserve the support of the prior. If $\nu = 0$, the posterior is zero at every history; if $\nu = 1$, it is one at every history. For $0 < \nu < 1$, an action by H that has zero probability under both types may support any posterior in $[0, 1]$. This is a condition of our solution concept. It follows the no-signaling logic used in multistage games (Fudenberg and Tirole 1991); support preservation is analogous to, but narrower than, the never-dissuaded discipline discussed by Osborne and Rubinstein (1990). Tremble consistency motivates the restriction but is not invoked as a theorem for degenerate priors (Kreps and Wilson 1982).

Table 2: Scope of the formal results

Element	Maintained scope
Players and horizon	One hegemon, $m \geq 3$ weak states, two rounds
Agenda	Only weak states propose in both rounds
Voting	Simultaneous public ballots; majority or unanimity
Information	Binary private type of H ; support-preserving endpoints
Payoffs	Unit pie, $0 < o_0 < o_1 < 1$, $0 < \beta < 1$, $b_\theta = 0$
Equilibrium	Perfect Bayesian equilibrium in pure ballot strategies with the declared voting and proposal tie-break rules

4 Results

The solution proceeds by backward induction. We begin with public types because that benchmark separates the value of a necessary vote from the value of private information. We then solve the private terminal games, the private first-round games, and finally subtract the public benchmark component by component.

4.1 Complete-information benchmark

With a public type, no belief or signaling issue remains. The proposer compares the price of H 's vote with the price of enough weak-state votes. In the terminal round, the relevant price of H is o_θ . In the first round, it is βo_θ , because rejection leads to the terminal round. Under majority, a

weak-state vote costs β/m in first-round units. Under unanimity with public disagreement payoff o , each weak responder's continuation price is $\beta(1 - o)/m$.

Proposition 4.1 (Public-type benchmark). *Fix a public type with disagreement payoff o . In the terminal round, majority passes a proposal without H , assigns $y = 0$, and gives the weak proposer the whole unit pie; H receives o . Terminal unanimity passes with $y = o$, and the weak proposer retains $1 - o$.*

In Round 1, unanimity passes immediately with $y = \beta o$. Majority includes H if $o \leq 1/m$ and excludes it if $o > 1/m$. At equality the proposal including H is selected. Consequently, the public Round-1 payoff of H is

$$p_M(o) = \begin{cases} \beta o, & o \leq 1/m, \\ o, & o > 1/m, \end{cases} \quad p_U(o) = \beta o.$$

Table 3: Complete-information outcomes by round and voting rule

Round	Rule	Proposed payments	Weak proposer	H
2	Majority	$y = 0$; weak responders get zero	1	o
2	Unanimity	$y = o$; weak responders get zero	$1 - o$	o
1	Majority, $o \leq 1/m$	$y = \beta o$; $q - 2$ responders get β/m	$1 - \beta o - (q - 2)\beta/m$	βo
1	Majority, $o > 1/m$	$y = 0$; $q - 1$ responders get β/m	$1 - (q - 1)\beta/m$	o
1	Unanimity	$y = \beta o$; every responder gets $\beta(1 - o)/m$	$1 - \beta o - (m - 1)\beta(1 - o)/m$	βo

At $o = 1/m$, buying H 's vote and buying the substitute weak-state vote cost the proposer the same amount. The proposal tie-break selects inclusion because it gives H the lower payoff $\beta o < o$.

4.2 Private information in the terminal round

Let

$$\nu^* = \frac{o_1 - o_0}{1 - o_0}.$$

Under terminal majority, H 's vote is unnecessary and the belief never enters the proposer's choice. Under terminal unanimity, the proposer compares a low offer that succeeds only when the type is low with a high offer that both types accept.

Proposition 4.2 (Private terminal games). *Under majority, the unique equilibrium outcome assigns $y = 0$, pays no responding weak state, and gives the full pie to the proposer. The proposal passes without H ; type θ receives o_θ , and a weak state receives $1/m$ before recognition.*

Under unanimity, if $\nu \leq \nu^$, the proposer offers $y = o_0$. The low type accepts and the high type rejects, so the proposal succeeds with probability $1 - \nu$. If $\nu > \nu^*$, the proposer offers $y = o_1$ and both types accept. At $\nu = \nu^*$, the low offer is selected.*

The cutoff equates $(1 - \nu)(1 - o_0)$, the proposer's payoff from the low offer, with $1 - o_1$, its payoff from pooling. Notice that no discount factor appears inside this terminal comparison.

4.3 Private majority in Round 1

Define

$$w = \frac{\beta}{m}, \quad t_\theta = \beta o_\theta,$$

and the proposer's payoffs from the three undominated outcome classes:

$$\begin{aligned} E &= 1 - (q - 1)w, \\ L &= 1 - (q - 2)w - t_0, \\ S(\nu) &= (1 - \nu)L + \nu w, \\ P &= 1 - (q - 2)w - t_1. \end{aligned}$$

Exclusion buys $q - 1$ weak responders at price w each and sets $y = 0$. Screening buys $q - 2$ weak responders and offers t_0 : the low type accepts and the high type delays. Pooling buys $q - 2$ weak responders and offers t_1 , which both types accept. Deliberate delay is strictly worse than exclusion

because $1 - \beta q/m > 0$.

When $o_0 < 1/m$, define

$$\nu_{SE} = \frac{\beta(1/m - o_0)}{\beta(1/m - o_0) + 1 - \beta q/m}.$$

When $o_1 < 1/m$, define

$$\nu_{SP} = \frac{\beta(o_1 - o_0)}{1 - \beta o_0 - \beta(q - 1)/m}.$$

Proposition 4.3 (Private majority correspondence). *For $m \geq 3$, the equilibrium outcome class is:*

1. *If $o_1 < 1/m$, screening for $\nu \leq \nu_{SP}$ and pooling above it.*
2. *If $o_0 < 1/m < o_1$, screening for $\nu \leq \nu_{SE}$ and exclusion above it.*
3. *If $1/m < o_0 < o_1$, exclusion for every ν .*
4. *If $o_0 = 1/m < o_1$, screening at $\nu = 0$ and exclusion for $\nu > 0$.*
5. *If $o_0 < o_1 = 1/m$, screening for $\nu \leq \nu_{SE}$. Above that cutoff, exclusion is selected when $(1 - \nu)o_0 + \nu o_1 < \beta o_1$, pooling is selected when the inequality is reversed, and the entire proposal segment joining the two survives when the two expressions are equal.*

At every cutoff involving screening, screening is selected. Permuting the identities of identically paid weak responders generates additional equilibria but does not change H 's payoff vector or the outcome class. Payoffs of weak states identified by label can be permuted and, on the residual proposal family, vary with the common proposal weight.

The last segment is genuine multiplicity. Its mixing weight is a probability over pure proposals in the proposer's strategy, and the same weight must be used jointly for payoffs and outcomes. It is not a license to combine componentwise minima and maxima into an unattained rectangle.

4.4 Private unanimity in Round 1

Let

$$\ell = \beta o_0, \quad h = \beta o_1, \quad A = \frac{\beta(1 - o_0)}{m}, \quad B = \frac{\beta(1 - o_1)}{m}.$$

At $\nu = 0$, support preservation prevents an impossible high type from being reintroduced by an off-path belief. The proposer can therefore offer the low threshold ℓ , pay each responding weak

state its continuation A , and keep $A + 1 - \beta$. When $\nu > \nu^*$, the continuation itself pools, so the proposer offers h , pays each responder B , and keeps $B + 1 - \beta$.

The intermediate cell fails for a different reason. A feasible deviation can pay every weak responder enough to force a yes vote. Once all weak votes are yes, no pure contingent voting profile for the two possible types of H survives: pooling yes is rejected by the high type, pooling no is overturned by the low type at equality, and either separating profile gives one type a profitable imitation.

Proposition 4.4 (Private unanimity correspondence). *Under unanimity:*

- $\nu = 0$: immediate low-type agreement with $y = \ell$, $x_j = A$, $r_i = A + 1 - \beta$;
- $0 < \nu \leq \nu^*$: no perfect Bayesian equilibrium in pure ballot strategies;
- $\nu^* < \nu \leq 1$: immediate pooling agreement with $y = h$, $x_j = B$, $r_i = B + 1 - \beta$.


The corresponding payoff vectors for H , ordered as low and high type, are (ℓ, h) at $\nu = 0$ and (h, h) above ν^* .

Remark (Pure-strategy scope). The middle cell establishes nonexistence only within the paper's maintained space of pure ballot strategies. We do not derive, characterize, select, or compare equilibria involving mixed ballot strategies, and make no claim about existence under such strategies.

Table 4: Private-information Round-1 equilibrium correspondences

Rule	Parameter cell	Selected outcome class	H payoff vector
Majority	$o_1 < 1/m$	Screening for $\nu \leq \nu_{SP}$; pooling above	(t_0, t_1) or (t_1, t_1)
Majority	$o_0 < 1/m < o_1$	Screening for $\nu \leq \nu_{SE}$; exclusion above	(t_0, t_1) or (o_0, o_1)
Majority	$1/m < o_0$	Exclusion for every ν	(o_0, o_1)
Majority	$o_0 = 1/m < o_1$	Screening at $\nu = 0$; exclusion for $\nu > 0$	(t_0, t_1) or (o_0, o_1)
Majority	$o_0 < o_1 = 1/m$	Screening through ν_{SE} ; above, exclusion, pooling, or their exact proposal segment under the tie-break	(t_0, t_1) , (o_0, o_1) , (t_1, t_1) , or the linked segment
Unanimity	$\nu = 0$	Immediate low-type agreement	(ℓ, h)
Unanimity	$0 < \nu \leq \nu^*$	No PBE in pure ballot strategies	$\emptyset$
Unanimity	$\nu^* < \nu \leq 1$	Immediate pooling agreement	(h, h)

Prices and coalition anatomy



Majority caps the hegemon's price by buying substitute votes; unanimity pays the hegemon directly. Panel A plots type-specific round-1 payoffs. Majority pays $\beta \times o_0$ and $\beta \times o_1$ through nu_{SE} , then exclusion leaves H with its outside option; unanimity has no pure-vote PBE for $0 < nu < nu^*$ and pools at $h = \beta \times o_1$ above nu^* . Under unanimity both types receive h ; the dashed high-type line is drawn over the solid low-type line where they coincide. The yellow span is the low type's pooling rent $h - \ell$, not the public-benchmark rent estimand. Panel B decomposes the unit surplus at $nu = 0.35$: majority buys $q-1$ substitutes at β/m , while the adjacent markers show the excluded H collecting o_{θ} outside the pie; unanimity pays $m-1$ weak-state floors, h , and the proposer residual. Parameters: $o_0 = 0.100$, $o_1 = 0.350$, $m = 4$, $\beta = 0.90$. Note: Model-generated regions; all boundaries closed-form.

Figure 2: The price and coalition anatomy of the private games. Majority can replace the hegemon's vote with an additional weak-state vote; unanimity cannot. Labels report first-round prices and preserve the distinct low-type, pooling, and exclusion outcomes.

4.5 Comparing the private games

The comparison below uses the same parameter point under both institutions and only cells in which both games have a pure-strategy equilibrium. Let

$$a_{\theta} = (1 - \beta)o_{\theta}, \quad d = \beta(o_1 - o_0), \quad k = \beta o_1 - o_0.$$

The terms a_θ are timing wedges: exclusion realizes the disagreement payoff now, whereas inclusion can price a continuation that is discounted in Round-1 units. They are distinct from the type gap d and the cross-date quantity k .


Proposition 4.5 (Private institutional payoff contrast). *Write $V_U^{priv} - V_M^{priv}$ as a low-type/high-type vector.*

1. *At $\nu = 0$, the vector is $(0, 0)$ if majority screens and $(-a_0, -a_1)$ if majority excludes.*
2. *For $0 < \nu \leq \nu^*$, the contrast is empty because private unanimity has no pure-strategy equilibrium.*
3. *For $\nu > \nu^*$, the vector is $(d, 0)$ if majority screens, $(0, 0)$ if majority pools, and $(k, -a_1)$ if majority excludes.*
4. *On the residual exclusion–pooling segment at $o_1 = 1/m$, the exact attainable set is $\{\lambda(k, -a_1) : \lambda \in [0, 1]\}$, restricted to the proposal weights that survive the tie-break. The same λ binds payoffs and outcomes.*

Thus unanimity does not have a single unconditional payoff effect. Its effect depends on which substitute coalition disciplines the majority proposer and on which type is evaluated. Figure 3 maps these exact classes without averaging types.

The private-rule payoff contrast is type-specific

Closed-form slice $\alpha_0 = 0.50 \times \alpha_1$; $m = 4$ weak states; $\beta = 0.90$



Colored regions compare the hegemon's private-information payoff under unanimity and majority at the same parameter point. The hatched region is empty because unanimity has no perfect Bayesian equilibrium in pure ballot strategies. The left-edge segment is the literal $\nu = 0$ endpoint. The horizontal threshold $m \times \alpha_1 = 1$ is the cost at which majority switches between buying the hegemon and buying an additional weak-state vote. No prior-weighted or ex ante image is displayed. Note: model-generated closed-form regions.

Figure 3: Private-game payoff contrast, unanimity minus majority, by hegemon type. Each cell is evaluated at the same parameter point. The hatched band marks parameter values for which private unanimity has no pure-strategy equilibrium and the contrast is therefore empty.

4.6 Informational rents and the difference of differences

For rule $g \in \{M, U\}$, define the type-contingent informational-rent correspondence

$$RI_g = V_g^{priv} - V_g^{pub},$$

where the public benchmark is evaluated at the same type and the same remaining parameters. The institutional informational-rent contrast is

$$\Delta RI = RI_U - RI_M.$$

These are differences between correspondences. If a source correspondence is empty, so is the difference. Under multiplicity, subtraction preserves the entire attainable set and its common proposal weight.

For compactness, call the public outside-option regions

$$\text{II} : o_1 \leq 1/m, \quad \text{IX} : o_0 \leq 1/m < o_1, \quad \text{XX} : 1/m < o_0.$$

The labels only record whether each public type is included by majority.

The rent comparison first asks what private information changes *within* each rule. Under unanimity, the public benchmark already buys the known type at its discounted threshold; private pooling therefore raises only the low type's payoff. Under majority, information matters only when the private equilibrium changes the public inclusion decision or the price paid to a type.

Proposition 4.6 (Informational rents by voting rule). *Under unanimity,*

$$RI_U = \begin{cases} \{(0, 0)\}, & \nu = 0, \\ \emptyset, & 0 < \nu \leq \nu^*, \\ \{(d, 0)\}, & \nu > \nu^*. \end{cases}$$

Under majority, the informational rent depends jointly on the public region and the private equilibrium class:

$$\begin{aligned} \text{II, screening} & : (0, 0), \\ \text{II, pooling} & : (d, 0), \\ \text{II, exclusion} & : (a_0, a_1), \\ \text{IX, screening} & : (0, -a_1), \\ \text{IX, exclusion} & : (a_0, 0), \\ \text{XX, exclusion} & : (0, 0). \end{aligned}$$

On the residual exclusion–pooling segment at $o_1 = 1/m$, the exact set is $\{\lambda(a_0, a_1) + (1 - \lambda)(d, 0) : \lambda \in [0, 1]\}$, restricted to the segment that survives the proposal tie-break.

The difference of differences then removes majority's informational rent from unanimity's. Screening can leave a positive low-type contrast in the high-belief cell because unanimity pools. Exclusion instead compares the high threshold discounted to the present with the low current disagreement payoff, so its low-type sign is governed by $k = \beta o_1 - o_0$.

Proposition 4.7 (Institutional informational-rent contrast). *At $\nu = 0$, $\Delta RI = (0, 0)$ in regions II and XX, while $\Delta RI = (0, a_1)$ in region IX; the high-type component there is off support but remains part of the type-contingent vector.*

For $0 < \nu \leq \nu^$, $\Delta RI = \emptyset$.*

For $\nu > \nu^$, the exact vectors are*

$$\begin{aligned}
\text{II, screening} & : (d, 0), \\
\text{II, pooling} & : (0, 0), \\
\text{II, exclusion} & : (k, -a_1), \\
\text{IX, screening} & : (d, a_1), \\
\text{IX, exclusion} & : (k, 0), \\
\text{XX, exclusion} & : (d, 0).
\end{aligned}$$

On the residual segment, the exact attainable set is $\{\lambda(k, -a_1) : \lambda \in [0, 1]\}$ over the surviving proposal weights.

Table 5: Exact informational-rent correspondences and institutional contrast

Object	Cell	Exact set
RI_M	II, screening	$\{(0, 0)\}$
RI_M	II, pooling	$\{(d, 0)\}$
RI_M	II, exclusion	$\{(a_0, a_1)\}$
RI_M	II, exclusion–pooling segment	$\{\lambda(a_0, a_1) + (1 - \lambda)(d, 0)\}$
RI_M	IX, screening	$\{(0, -a_1)\}$
RI_M	IX, exclusion	$\{(a_0, 0)\}$
RI_M	XX, exclusion	$\{(0, 0)\}$
RI_U	$\nu = 0$	$\{(0, 0)\}$
RI_U	$0 < \nu \leq \nu^*$	$\emptyset$
RI_U	$\nu > \nu^*$	$\{(d, 0)\}$
ΔRI	$\nu = 0$, II	$\{(0, 0)\}$

Object	Cell	Exact set
ΔRI	$\nu = 0, IX$	$\{(0, a_1)\}$
ΔRI	$\nu = 0, XX$	$\{(0, 0)\}$
ΔRI	$0 < \nu \leq \nu^*$, every public region	$\emptyset$
ΔRI	$\nu > \nu^*$, II screening	$\{(d, 0)\}$
ΔRI	$\nu > \nu^*$, II pooling	$\{(0, 0)\}$
ΔRI	$\nu > \nu^*$, II exclusion	$\{(k, -a_1)\}$
ΔRI	$\nu > \nu^*$, II exclusion–pooling segment	$\{\lambda(k, -a_1) : \lambda \in [0, 1]\}$
ΔRI	$\nu > \nu^*$, IX screening	$\{(d, a_1)\}$
ΔRI	$\nu > \nu^*$, IX exclusion	$\{(k, 0)\}$
ΔRI	$\nu > \nu^*$, XX exclusion	$\{(d, 0)\}$

In every segment row, $\lambda \in [0, 1]$ is restricted to the proposal weights that survive the tie-break and binds all payoff and outcome components.

The sign pattern is transparent once each component is kept separate. Above the unanimity cut-off, the low type gains $d > 0$ when majority screens because unanimity pools and pays the high threshold. The same screening comparison is zero at $\nu = 0$, while the middle cell is empty. When majority excludes, the low type's gain is k , which is positive, zero, or negative according as βo_1 is above, equal to, or below o_0 . The high type gains $a_1 > 0$ only when private majority screens even though public majority would exclude it; it loses a_1 when public majority includes it but private majority excludes. When both private rules pool, the informational-rent contrast is exactly zero.

Power versus information




Figure 4: Power and information in the hegemon’s payoff. The public benchmark captures the effect of making the hegemon’s vote necessary when its type is known. Informational rent compares the private game with that benchmark, and ΔRI compares this rent across unanimity and majority without recombining types or selecting among multiple equilibria.

5 Discussion


5.1 Pivotality, substitutes, and contested strength

The model’s central distinction is technological. Under majority, the proposer can replace the hegemon with another uninformed weak-state vote. Under unanimity, the hegemon’s approval is an essential input. Removing the substitute changes both bargaining power and the way private information enters proposals.

The complete-information benchmark shows the power component. If the known hegemon is cheap enough, majority includes it and both rules pay the same discounted threshold. If it is expensive, majority excludes it while unanimity must still buy its vote. The informational comparison asks a different question: after netting out that public-type difference, how much does private information change each rule's payoff?

The answer is not simply that private information always benefits a hegemon. Pooling under unanimity benefits the low type because weak states price it as high. Screening under majority can remove that rent, while exclusion can make the comparison depend on whether the high threshold discounted to the present exceeds the low current disagreement payoff. The high type's rent changes only when its treatment differs across the private and public majority games.

The empty middle-belief cell is equally important for scope. It records the failure of every pure voting pattern under the declared belief and pivotality conditions. It does not supply a payoff for institutional comparison and does not justify filling the gap by interpolation.



The figure should be read from high to low ν . Above ν^* , pooling pays the low type $h = \beta \alpha_1$, creating the rent $h - \ell$ relative to its reservation price $\ell = \beta \alpha_0$. The light orange field identifies the pooling region; the darker band isolates the payoff-relevant rent between ℓ and h . For $0 < \nu < \nu^*$, the hatched area records the absence of a perfect Bayesian equilibrium in pure ballot strategies; $\nu = 0$ is an isolated complete-information endpoint, not an interval, and the low type is bought at reservation. The figure makes no claim about equilibria involving mixed ballot strategies. Note: Model-generated regions; all boundaries closed-form.

Figure 5: A schematic movement in beliefs about hegemonic strength. The hatched region is the set in which the unanimity game has no equilibrium in pure ballot strategies under the maintained concept. Historical annotations, if used, are illustrations rather than empirical tests or calibration.

5.2 OPEC, the WTO, and observable implications

In the OPEC interpretation, a proposed package can represent quota flexibility, an exception, or enforcement discretion. Saudi Arabia may know more than other members about the opportunity cost of accepting it. The model predicts that making its approval necessary removes a coalition substitute and can induce an offer calibrated to the higher possible reservation value. This is a mechanism illustration, not an estimate of OPEC bargaining.

For the WTO, consensus similarly makes a major state's vote necessary despite formally equal voting rights and the absence of exclusive proposal power. The model suggests looking for institutional concessions that arise because partners cannot confidently price the major state's willingness to accept. The relevant observable contrast is not merely whether the hegemon is included, but whether majority would provide a cheaper coalition of uninformed votes.

These implications differ from informational voting models in which individual votes aggregate dispersed signals (Feddersen and Pesendorfer 1998). Here weak states do not possess private signals about H 's type. They respond to a public proposal and to the continuation induced by the publicly revealed vote vector. The nearby legislative-bargaining literature analyzes additional equilibrium concepts and strategic forms (Piazolo and Vanberg 2025; Glynia, Thum, and Xefteris 2026); our comparison deliberately remains within the pure ballot strategies stated above.

5.3 Limits

The analysis fixes the voting rule, keeps the institutional pie at one, gives proposal power only to weak states, and uses two types and two rounds. It does not analyze endogenous rule choice, agenda power for H , more than two types, sequential roll-call voting, or intrinsic agreement benefits. The restriction $m \geq 3$ excludes the three-player case. The strict discounting assumption excludes the corner $\beta = 1$, where additional indifferences would require a separate analysis.

The endpoint-support condition also matters. A type assigned zero probability by the prior cannot be resurrected after an off-path action. This makes the degenerate-prior cells coincide with their complete-information games. In the interior, both types remain in the prior's support and off-path actions by H can sustain a range of posterior beliefs.

6 Conclusion

Consensus can benefit a hegemon even when weak states control the agenda because unanimity removes substitutes for the hegemon's vote. Complete information reveals the familiar veto-power component. Private information adds a distinct component: when unanimity pools at the high threshold, a low type can earn rent; whether this exceeds its majority rent depends on the majority equilibrium and the outside-option region.

The decomposition disciplines the claim. The institutional informational-rent contrast is positive in some regions, zero when majority already pays the pooling price, sign-indeterminate through $k = \beta o_1 - o_0$ under exclusion, and empty where unanimity has no pure-strategy equilibrium. Formal equality of votes can therefore coexist with asymmetric informational power, but the effect is neither universal nor separable from the coalition alternatives created by the voting rule.

Appendix A: Protocol, beliefs, and incentives

A.1 Complete transition and payoff rules

For a proposal $s = (y, (x_j), r_i)$, let $a_H \in \{Y, N\}$ and let the weak responders' votes be a_{-i} . The proposer votes yes. If the quota is met, the proposal is implemented in full. Each weak respondent receives x_j , the proposer receives r_i , and H receives y after yes and $y + o_\theta$ after no. If the quota is not met in Round 1, the public vote vector leads to Round 2; every continuation payoff is then multiplied by β . If the quota is not met in Round 2, weak states receive zero and H receives o_θ .

The $y + o_\theta$ branch is necessary for a complete strategy specification. It is not an equilibrium description. Whenever majority passes without H , a proposer strictly prefers to move any positive y into its own residual, holding all votes fixed. Hence all equilibrium exclusions set $y = 0$ and pay H exactly o_θ .

A.2 Beliefs and ballot restrictions

Uninformed weak-state proposals and votes preserve the current belief. A vote by H updates by Bayes' rule when its denominator is positive. At a zero-denominator action in an interior-prior game, the posterior may range over $[0, 1]$, subject to structural consistency across histories that encode the same information. At $\nu = 0$ or $\nu = 1$, posterior support never expands.

A weak responder compares its allocation with the continuation that would follow if its vote switched passage into failure. It votes yes at equality. H maximizes its type-contingent continuation payoff and votes yes at exact indifference. These restrictions apply at every proposal, including off-path proposals.

Appendix B: Proofs

B.1 Proof of Proposition 4.1

In terminal majority, a weak responder's disagreement value is zero, so any nonnegative allocation induces yes under T^Y . Because $q \leq m$, the proposer can pass without H , set every payment

to zero, and keep one. H is nonpivotal and strictly votes no because no yields $y + o$ rather than y . Moving any positive y to the proposer leaves passage unchanged, so equilibrium exclusion has $y = 0$ and H receives o .

In terminal unanimity, every weak responder accepts zero and H accepts exactly when $y \geq o$. The proposer therefore sets $y = o$.

Move to Round 1. The terminal continuation of a weak state before recognition under majority is $1/m$, hence $w = \beta/m$. The continuation of H is βo . Under unanimity, paying H exactly βo and each weak responder its continuation implements immediate agreement and leaves the proposer strictly more than delay because $\beta < 1$.

Under majority, inclusion costs $(q - 2)w + \beta o$; exclusion costs $(q - 1)w$. Inclusion is weakly cheaper exactly when $o \leq 1/m$. At equality the proposer is indifferent, and the proposal tie-break selects inclusion because H receives $\beta o < o$. This gives the stated payoffs. $\square$

B.2 Proof of Proposition 4.2

The majority argument is the terminal-majority part of the preceding proof. Before recognition, symmetry gives each weak state expected payoff $1/m$.

Under unanimity, weak responders accept zero. Any accepted offer to H can be reduced to the lowest threshold consistent with its acceptance set. The only nonempty acceptance sets are the low type alone, implemented by $y = o_0$, and both types, implemented by $y = o_1$. The first gives the proposer $(1 - \nu)(1 - o_0)$; the second gives $1 - o_1$. The low offer is weakly preferred exactly when

$$\nu \leq \frac{o_1 - o_0}{1 - o_0} = \nu^*.$$

At equality the low offer gives H the lower expected payoff and is selected. $\square$

B.3 Proof of Proposition 4.3

A responding weak state votes yes exactly when $x_j \geq w = \beta/m$. Let k be the number of such responders. If $k \geq q - 1$, the weak votes pass the proposal without H ; both types of H vote no

and receive $y + o_\theta$. If $k = q - 2$, H is pivotal and type θ votes yes exactly when $y \geq t_\theta = \beta o_\theta$. If $k \leq q - 3$, the proposal fails regardless of H 's vote.

Within each acceptance class, reducing weak payments to w , reducing y to the relevant type threshold, and assigning the remaining pie to the proposer weakly increases its payoff. This leaves four candidates: exclusion with payoff E , screening with payoff $S(\nu)$, pooling with payoff P , and deliberate delay with payoff w . Exclusion is always feasible and

$$E - w = 1 - \frac{\beta q}{m} > 0,$$

so delay is never selected. A screening or pooling candidate that could beat exclusion is automatically feasible.

The decisive payoff differences are

$$P - E = \beta(1/m - o_1)$$

and

$$S(\nu) - E = (1 - \nu)\beta(1/m - o_0) - \nu(1 - \beta q/m).$$

Where pooling can beat exclusion, comparing screening with pooling yields ν_{SP} ; comparing screening with exclusion yields ν_{SE} . Substituting the signs of $o_0 - 1/m$ and $o_1 - 1/m$ produces the five cases in Proposition 4.3.

At screening ties, its expected payoff to H is strictly smaller than that of the tied alternative, so the proposal tie-break selects screening. When $o_1 = 1/m$, exclusion and pooling give the proposer the same payoff. Their payoffs to H determine the stated selection; exact equality leaves their entire proposal segment. All selected proposals exhaust the pie. Finally, permuting the identities of weak responders does not change any comparison, which establishes the reported identity multiplicity. $\square$

B.4 Proof of Proposition 4.4

Transport the terminal unanimity continuation into Round-1 units. If the posterior is η , a weak state's continuation is

$$W(\eta) = \begin{cases} (1 - \eta)A, & \eta \leq \nu^*, \\ B, & \eta > \nu^*. \end{cases}$$

Thus $W(\eta) \in [B, A]$, and $x_j = A$ forces every weak responder to vote yes under any admissible posterior.

At $\nu = 0$, support preservation fixes every posterior at zero. The proposal

$$y = \ell, \quad x_j = A, \quad r_i = 1 - \ell - (m - 1)A = A + 1 - \beta$$

is accepted. Each responder and H is exactly indifferent and votes yes. Any lower payment loses a required vote; any higher payment reduces the proposer's residual. A delay gives the proposer A , so immediate agreement is strictly better by $1 - \beta$.

For $\nu > \nu^*$, the analogous pooling proposal is

$$y = h, \quad x_j = B, \quad r_i = 1 - h - (m - 1)B = B + 1 - \beta.$$

Both types of H and every weak responder accept, and immediate agreement again exceeds delay by $1 - \beta$.

For completeness, write $u = \min_j x_j$ after an arbitrary proposal and order the strategy of H as low-type/high-type. At $\nu = 0$, posterior support is always the low type. The profile (Y, Y) is sequentially rational if $u < A$, or if $u \geq A$ and $y \geq h$; (N, N) is rational if $u \geq A$ and $y < \ell$; and (Y, N) is rational if $u \geq A$ and $\ell \leq y < h$. The profile (N, Y) is never rational. These cases are mutually exclusive and exhaustive.

For $\nu^* < \nu < 1$, the profile (Y, Y) is rational if $u < B$, or if $u \geq B$ and $y \geq h$; (N, N) is rational if $u \geq B$ and $y < h$; and neither separating profile is rational. Bayes gives the current prior after an on-profile pooled action. When yes is a zero-probability action, an admissible posterior η_Y supports all weak yes votes exactly when $W(\eta_Y) \leq u$, which is possible in the stated no profile because

$u \geq B$. At $\nu = 1$, support preservation fixes the posterior at one and the same two pooled-profile conditions apply. Thus every proposal in both existence domains has a pure, sequentially rational response completion.

It remains to show the empty cell. For $0 < \nu \leq \nu^*$, consider

$$s^\dagger : \quad y = \ell, \quad x_j = A \text{ for every weak responder,} \quad r_i = 1 - \ell - (m - 1)A.$$

All weak responders must vote yes. Enumerate the four pure type-contingent profiles for H , ordered low/high:

1. Under (Y, Y) , the high type receives $\ell < h$ and profitably votes no to obtain continuation h .
2. Under (N, N) , the low type's no gives ℓ , so T^Y requires yes at equality.
3. Under (Y, N) , the low type can imitate the high type's no. That action reveals the high type and yields continuation $h > \ell$.
4. Under (N, Y) , the high type can imitate the low type's no and obtain continuation h , while the prescribed comparison cannot support strict no against yes at the offer.

No pure profile is sequentially rational after $s^\dagger$. Because a PBE must specify a valid pure response after every feasible proposal, no pure-ballot PBE exists in this cell. This also explains the discontinuity: at $\nu = 0$, support preservation rules out the posterior that drives the imitation argument; for every positive ν , both types are possible. $\square$

B.5 Proof of Proposition 4.5

The private majority payoff vectors for H are

$$V_M^S = (\beta o_0, \beta o_1), \quad V_M^P = (\beta o_1, \beta o_1), \quad V_M^E = (o_0, o_1).$$

The private unanimity vectors are $(\beta o_0, \beta o_1)$ at $\nu = 0$, empty for $0 < \nu \leq \nu^*$, and $(\beta o_1, \beta o_1)$ for $\nu > \nu^*$. Componentwise subtraction gives the proposition. On any proposal segment, subtraction is affine in the common proposal weight, so the payoff set and outcome set remain atomic. $\square$

B.6 Proofs of Propositions 4.6 and 4.7

From Proposition 4.1, the public payoff vectors are

$$V_M^{pub} = (p_M(o_0), p_M(o_1)), \quad V_U^{pub} = (\beta o_0, \beta o_1).$$

Subtracting these from each private vector gives the table in Proposition 4.6. In particular, unanimity is payoff-equivalent to its public endpoint at $\nu = 0$, is empty in the middle cell, and pools above the cutoff, which yields $(d, 0)$.

Subtract RI_M from RI_U . In the high-belief cell, screening gives $(d, 0)$ in region II and (d, a_1) in region IX; pooling gives zero; exclusion gives $(k, -a_1)$, $(k, 0)$, or $(d, 0)$ in regions II, IX, or XX. The endpoint follows by the same calculation. An empty source set remains empty. For the residual segment, affine subtraction with its common weight λ yields the one-dimensional segment stated in the propositions, not the Cartesian product of marginal envelopes. $\square$

Appendix C: Endpoints, exact sets, and illustration

C.1 Endpoint equivalence

At $\nu = 0$, posterior support is the singleton low type. The private unanimity proposal and payoff coincide with the public low-type game. At $\nu = 1$, posterior support is the singleton high type, and the private pooling proposal and payoff coincide with the public high-type game. The same equivalence holds under majority: applying the relevant private class at each endpoint reproduces public inclusion or exclusion, including the $o_\theta = 1/m$ tie-break. Thus the endpoints are literal complete-information games, not one-sided limits.

C.2 Exact correspondence and envelopes

Most parameter cells are singleton payoff vectors. At $o_1 = 1/m$, exact indifference between exclusion and pooling can leave a proposal segment. If λ is the exclusion weight, every reported object uses that same λ . The lower and upper envelope of any component may summarize the segment, but coordinated marginal envelopes do not imply that all pairwise combinations are attainable.

For majority-rule informational rents, the exact componentwise envelopes are

$$RI_M(0) \in [\min\{a_0, d\}, \max\{a_0, d\}], \quad RI_M(1) \in [0, a_1].$$

For both the private institutional payoff contrast and ΔRI , they are

$$\text{low type} \in [\min\{0, k\}, \max\{0, k\}], \quad \text{high type} \in [-a_1, 0].$$

In each case these are the marginal envelopes of a single line segment. The same λ indexes both coordinates, so the Cartesian product of the two intervals is not attainable.

Where $0 < \nu \leq \nu^*$, the private unanimity correspondence is empty. Accordingly, V_U^{priv} , RI_U , and ΔRI are empty. No payoff, sign, institutional ranking, or interpolation is assigned to that cell.

C.3 Worked values

Return to $m = 4$, $q = 3$, $\beta = 0.9$, $o_0 = 0.10$, and $o_1 = 0.35$. Besides $\nu^* = 0.2778$, the majority screening–exclusion cutoff is

$$\nu_{SE} = \frac{0.9(0.25 - 0.10)}{0.9(0.25 - 0.10) + 1 - 0.9(3/4)} = 0.2935.$$

At $\nu = 0.80$, majority excludes and unanimity pools. The public payoff vectors are

$$V_M^{pub} = (0.09, 0.35), \quad V_U^{pub} = (0.09, 0.315),$$

while the private payoff vectors are

$$V_M^{priv} = (0.10, 0.35), \quad V_U^{priv} = (0.315, 0.315).$$

Therefore

$$RI_M = (0.01, 0), \quad RI_U = (0.225, 0), \quad \Delta RI = (0.215, 0).$$

These numbers illustrate one point in region IX; they do not average over parameter values or estimate a real institution.

Appendix D: Notation

Table 6: Notation for the essential-input bargaining game

Symbol	Meaning
H	Privately informed hegemon
W, m	Set and number of weak states; $m \geq 3$
$N = m + 1$	Total number of states
q	Majority quota, $\lfloor N/2 \rfloor + 1$
θ	Hegemon type, low 0 or high 1
ν	Prior probability of the high type
o_θ	Type- θ terminal disagreement payoff
β	One-round discount factor, strictly between zero and one
y	Institutional concession assigned to H
$\bar{y}$	Maximum feasible institutional concession; $o_1 \leq \bar{y} \leq 1$
x_j	Allocation assigned to weak responder j
r_i	Residual assigned to weak proposer i
$w = \beta/m$	First-round continuation value of a weak state under majority
$t_\theta = \beta o_\theta$	First-round threshold of type θ when pivotal
ν^*	Terminal unanimity cutoff, $(o_1 - o_0)/(1 - o_0)$
a_θ	Timing wedge, $(1 - \beta)o_\theta$
d	Discounted type gap, $\beta(o_1 - o_0)$
k	Cross-date quantity, $\beta o_1 - o_0$
V_g^{priv}, V_g^{pub}	Private- and public-information payoff correspondences under rule g
RI_g	Informational-rent correspondence $V_g^{priv} - V_g^{pub}$
ΔRI	Unanimity-minus-majority informational-rent contrast

References

- Baron, David P., and John A. Ferejohn. 1989. "Bargaining in Legislatures." *American Political Science Review* 83 (4): 1181–1206.
- Eraslan, Hülya, and Kirill S. Evdokimov. 2019. "Legislative and Multilateral Bargaining." *Annual Review of Economics* 11: 443–72.
- Feddersen, Timothy J., and Wolfgang Pesendorfer. 1998. "Convicting the Innocent: The Inferiority of Unanimous Jury Verdicts Under Strategic Voting." *American Political Science Review* 92 (1): 23–35. <https://doi.org/10.2307/2585926>.
- Fudenberg, Drew, and Jean Tirole. 1991. *Game Theory*. Cambridge, MA: MIT Press.
- Glynia, Nektaria, Marcel Thum, and Dimitrios Xefferis. 2026. "Unanimity Versus Majority: Proposing Under Incomplete Information." *Public Choice*. <https://doi.org/10.1007/s11127-026-01383-9>.
- Kalandrakis, Tasos. 2006. "Proposal Rights and Political Power." *American Journal of Political Science* 50 (2): 441–48.
- Koremenos, Barbara, Charles Lipson, and Duncan Snidal. 2001. "The Rational Design of International Institutions." *International Organization* 55 (4): 761–99.
- Krasner, Stephen D. 1983. "Structural Causes and Regime Consequences: Regimes as Intervening Variables." In *International Regimes*, edited by Stephen D. Krasner, 1–21. Ithaca, NY: Cornell University Press.
- Kreps, David M., and Robert Wilson. 1982. "Sequential Equilibria." *Econometrica* 50 (4): 863–94.
- Miller, Luis, Maria Montero, and Christoph Vanberg. 2018. "Legislative Bargaining with Heterogeneous Disagreement Values: Theory and Experiments." *Games and Economic Behavior* 107: 60–92. <https://doi.org/10.1016/j.geb.2017.11.003>.
- Osborne, Martin J., and Ariel Rubinstein. 1990. *Bargaining and Markets*. San Diego: Academic Press.
- Piazolo, David, and Christoph Vanberg. 2025. "Legislative Bargaining with Private Information: A Comparison of Majority and Unanimity Rule." *Games and Economic Behavior* 153: 499–522. <https://doi.org/10.1016/j.geb.2025.07.010>.
- Steinberg, Richard H. 2002. "In the Shadow of Law or Power? Consensus-Based Bargaining and

- Outcomes in the GATT/WTO.” *International Organization* 56 (2): 339–74.
- Stone, Randall W. 2011. *Controlling Institutions: International Organizations and the Global Economy*. Cambridge: Cambridge University Press.
- Winter, Eyal. 1996. “Voting and Vetoing.” *American Political Science Review* 90 (4): 813–23.
<https://doi.org/10.2307/2945844>.